CORTEX™

EXECUTIVE AI READINESS BRIEF

Generated: October 7, 2025

ID: 63D0967C

COMPLETED

EXECUTIVE SUMMARY

Overall AI Readiness: Integrated (1.6/3)

This assessment evaluates your organization across 6 CORTEX domains. 4 require focused development. 2 are in early stages.

Domain Performance

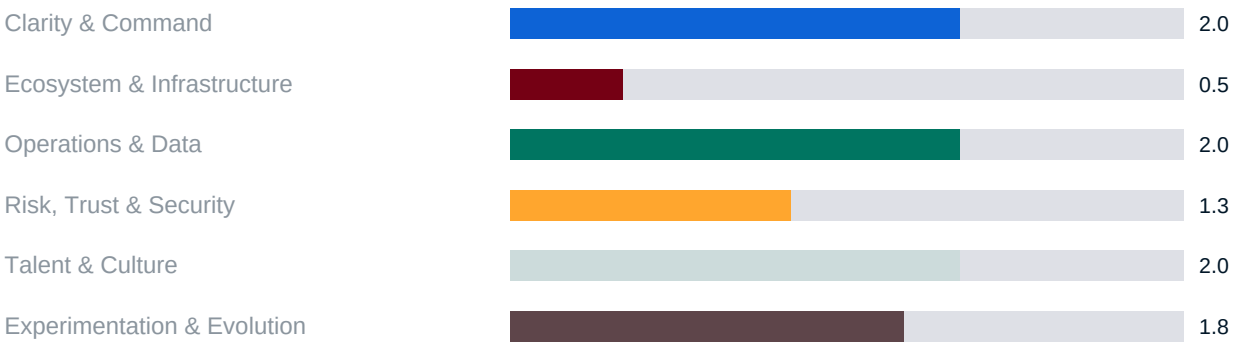

ORGANIZATIONAL CONTEXT

Your assessment captured the following organizational dimensions that shape your AI readiness requirements:

- Market Clock Speed: Monthly (2/4)
- Latency Requirements: <1s (1/4)
- Brand Exposure: Meaningful Risk (2/4)
- Data Advantage: Moderate (2/4)
- Build Readiness: Basics in Place (2/4)
- Edge Operations: No
- FinOps Priority: Medium (2/4)
- Data Sensitivity: Confidential (2/4)
- Scale/Throughput: Enterprise (2/4)
- Safety Criticality: Costly Mistakes (2/4)
- Regulatory Intensity: Some Rules (2/4)
- Procurement Requirements: No

ACTION PRIORITIES

1. Systematize AI Development Practices

Timeframe: 90 days

Establish governance, standardize processes, expand pilots

2. Strengthen Ecosystem & Infrastructure

Timeframe: 90 days

Focus improvement efforts on your biggest capability gap

DOMAIN ANALYSIS

C. Clarity & Command

Integrated (2.0/3)

Why This Matters

Clarity turns AI from scattered pilots into business outcomes. When leaders publish a simple, measurable AI ambition and name a single owner with budget authority, teams stop guessing. A routine executive review creates momentum: work that moves the needle is funded and scaled; experiments that don't deliver are sunset. This alignment reduces duplicated effort, accelerates learning, and ties AI to revenue, cost, and risk.

What Good Looks Like

- A one page AI ambition linked to business outcomes and customers
- A named senior owner and a clear split between CoE (enable/govern) and BUs (adopt/deliver)
- Quarterly executive/board review with reallocation decisions (fund/defund)
- Leaders share a common language for AI scope, risks, and value

How to Improve

Progress usually starts with publishing a simple ambition (outcomes, not technologies), then clarifying who owns what between a Center of Excellence and business units. Reviews move from "show and tell" to decide and do—small amounts of money shift to what w
Over time, AI outcomes appear in strategy documents and operating plans. In more regulated settings, leadership reviews also check that safeguards and evidence are in place.

Common Pitfalls to Avoid

- Vision without ownership; ownership without budget
- Treating the CoE as a gatekeeper instead of an enabler
- Endless exploration with no reallocation

Strategic Discussion Questions

- What two business outcomes will AI influence this year?
- Who is accountable for those outcomes and what budget do they contro
- What will we stop doing if it doesn't perform?

O. Operations & Data

Integrated (2.0/3)

Why This Matters

Stable operations and governed data are the difference between a demo and a dependable service. Monitoring, human in the loop where risk warrants it, and basic data hygi surprise bills, and reputational harm.

What Good Looks Like

- A documented lifecycle: design !' deploy !' monitor !' update !' retire
- Logging, alerts, and simple dashboards for usage, cost, latency, and failure rates
- Human review and quality assurance checkpoints where stakes are high
- A searchable data catalogue with owners, lineage, quality thresholds
- A lightweight value/feasibility gate for new use cases

How to Improve

Start with monitoring what you already run (latency, cost, error rate) and add simple alerts. Introduce a two page intake for new ideas: value hypothesis, data sources, risk level. key tables or content used by AI. Where decisions affect customers, add human approval until you have evidence that automation is safe.

Common Pitfalls to Avoid

- Over engineering MLOps before any value has shipped
- No drift or cost alerts; discovering issues from users or invoices
- Unowned data; stale or inconsistent sources

Strategic Discussion Questions

!' What do we measure today on our AI services? What's missing?

!' Which one dataset, if cleaned and owned, would unlock the most value

!' Where should a human stay in the loop for now?

R. Risk, Trust & Security

Emerging (1.3/3)

Why This Matters

Trust and safety make AI adoption sustainable. Stakeholders expect you to know what AI you run, the risks it carries, and how you'll respond when something goes wrong. Basic assurance practices prevent reputational damage and regulatory setbacks.

What Good Looks Like

- An AI inventory with owners and risk levels
- Scheduled checks for fairness, privacy, and model/data drift
- Periodic red teaming for prompts/jailbreaks and data exfiltration attem
- An incident response runbook with roles and communications
- Internal or third party review of controls (as required)

How to Improve

Catalog what you already use (systems, vendors, purpose, data). Schedule basic checks for high impact use cases and test your defenses with simple adversarial pro plan: who triages, who decides, who informs customers. Regulated contexts often add annual assurance whether internal or external.

Common Pitfalls to Avoid

- Policy documents without monitoring
- Unknown owners; no one reacts when metrics drift
- Treating red teaming as a one time event

Strategic Discussion Questions

- Which AI system could create the most damage if it failed? Do we mon
- Who picks up the phone when an AI incident occurs?
- How often do we test for bias, privacy, and jailbreaks?

T. Talent & Culture

Integrated (2.0/3)

Why This Matters

Adoption is about work, not tools. People need role specific skills and up when to use AI, when to verify, and how to escalate. Stories and incentives help good behaviors spread.

What Good Looks Like

- Clear job families with role based AI fluency
- SOPs/SOP checklists updated to include AI tasks and checkpoints
- "Wins and lessons" shared on a regular rhythm
- Incentives that reward safe, effective use

How to Improve

Pick two or three job families that touch customers or costly processes. Create before/after task maps and add simple guardrails (checklists, approval steps). Offer short, role s examples. Share what works and what fails—both teach.

Common Pitfalls to Avoid

- Generic training without job redesign
- Incentives that reward activity over outcomes
- "One wizard" knows everything; no diffusion

Strategic Discussion Questions

- Which roles will benefit most from AI in 90 days?
- What checkpoints keep customers safe while we learn?
- How will we recognize and reward smart usage?

E. Ecosystem & Infrastructure

Nascent (0.5/3)

Why This Matters

Partners and platform choices determine speed, cost, and flexibility. Elastic capacity keeps teams moving; portability and clear terms help you avoid lock in and surprises.

What Good Looks Like

- Elastic capacity and simple FinOps visibility (unit costs, quotas)
- Strategic partners that fill capability gaps
- Exit/portability plans in contracts (export formats, second source)
- Governed APIs and basic interoperability standards

How to Improve

Start by measuring unit costs and watching quotas. Consolidate on a few clear terms ("no training on our data/outputs" when needed). Draft a one switch, what we'd export, and a secondary option for critical paths.

Common Pitfalls to Avoid

- Vendor lock in via proprietary formats and unclear rights
- Quota bottlenecks; budget surprises
- One off integrations that don't scale

Strategic Discussion Questions

- !’ Which costs or quotas block us most often?
- !’ What contractual term would protect our data and options?
- !’ If our primary vendor failed tomorrow, what's our plan?

X. Experimentation & Evolution

Integrated (1.8/3)

Why This Matters

AI changes quickly. Disciplined experimentation—safe sandboxes, small budgets, explicit success and sunset criteria—increases learning velocity and prevents "pilot purgatory."

What Good Looks Like

- A guarded sandbox with representative data and spending caps
- A ring fenced slice of time/credits for experiments
- Every pilot has success and sunset criteria and a decision date
- Lightweight horizon scanning of tech, policy, and competitors

How to Improve

Provide a clear on ramp: where to try ideas, what's allowed, and how to r metric and decision date for every pilot. Run a short horizon brief quarterly to decide what to watch or ignore. Retire experiments on time so resources return to the pool.

Common Pitfalls to Avoid

- Pilots with no metrics or end dates
- Sandboxes with real data but no guardrails

- Chasing every new model without a hypothesis

Strategic Discussion Questions

- Which experiments have been running for months without a decision?
- What would "good enough" look like to promote or retire a pilot?
- Who scans the horizon, and what did they flag this quarter?

STRATEGIC INSIGHTS

Systematize AI Development Practices

You have foundational capabilities but lack systematic approaches to scale effectively.